
A Spiking Ring-Attractor Architecture for Representation and Control of Revolute Robot Joints

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Robotic systems with revolute joints evolve on circular state spaces, yet most
2 control architectures represent these states using Euclidean variables that ignore
3 this underlying geometry. We propose a fully spiking neural control architecture in
4 which ring attractors encode joint states, enforcing representations consistent with
5 the topology of revolute systems.

6 Homeostatic regulation computes goal-state mismatch directly over these dis-
7 tributed representations, and opponent gain-modulated populations convert this
8 mismatch into controlled phase shifts of the activity bump, implementing set-point
9 regulation as motion on the ring.

10 We show that the bump dynamics follow an $SO(2)$ geometry, and that gain asym-
11 metries induce a signed phase velocity on the ring. To interface distributed activity
12 with scalar regulation signals, we derive an analytic Fourier phase decoder and
13 implement its signed components using push-pull spiking populations.

14 For multi-joint systems, we demonstrate that local ring-based Fourier features
15 compose into product representations that linearly decode $SO(3)$ orientations
16 through forward kinematics. Experiments in NEST validate the full $SO(2)$ regula-
17 tion mechanism and the spiking product-feature decoder for lift, pitch, and yaw.
18 Robot experiments on a Tiago arm show that decoded ring dynamics generate joint
19 velocity commands, producing directed motion with joint-dependent accuracy.

20 Together, these results establish a neuromorphic control framework in which repre-
21 sentation, regulation, and action are unified within a ring-attractor substrate aligned
22 with the geometry of revolute systems.

23 1 Introduction

24 Neural circuits in biological systems often do not appear to separate representation from control.
25 Internal states are maintained as distributed activity patterns, goals can be encoded within the same
26 substrate, and behavior can emerge through continuous regulation of one activity state toward another.
27 This organization is especially clear in neural systems for orientation and navigation, where heading,
28 goal direction, and steering commands are represented and transformed within population activity
29 [7, 20, 13]. Ring attractors provide a particularly relevant example: in biological head-direction
30 systems, they maintain localized activity bumps on a circular manifold and support angular integration
31 [17, 18, 10]. This circular structure has also made ring attractors useful in computational and robotic
32 models of path integration, grid-cell organization, and bio-inspired navigation [11, 19, 15, 2].

33 Building on this circular population structure, we propose a robotic control architecture in which
34 ring attractors represent revolute joint states and serve as dynamical substrates for control. The
35 central idea is to regulate the phase of a persistent activity bump directly, rather than decoding joint
36 state into an external Euclidean variable before control. Homeostatic comparison provides a signed

37 mismatch between goal and state, while opponent gain modulation converts that mismatch into
 38 directed transport of the bump around the ring [3, 16, 22]. Because the neural activity bump and
 39 the physical joint angle both evolve on circular manifolds, this yields a topology-preserving, fully
 40 spiking control primitive for revolute variables.

41 This perspective also provides a natural geometric and compositional structure. Phase shifts of a
 42 ring-attractor bump correspond to motion on $SO(2)$, with infinitesimal displacements described by a
 43 Lie-algebraic tangent direction [12, 21, 9]. To interface this distributed representation with scalar
 44 regulatory signals, we derive an NEF-inspired Fourier decoder that extracts phase from ring activity
 45 [4]. For multi-joint systems, the same local Fourier features, $\cos q_i$ and $\sin q_i$, compose through
 46 known forward kinematics to provide an $SO(3)$ orientation readout. Thus, individual rings generate
 47 controlled $SO(2)$ trajectories, while their decoded features support robot-level orientation estimates.

48 Our contributions are: (i) a biologically inspired closed-loop ring-attractor controller in which a
 49 Hammel-style homeostatic comparator drives opponent gain modulation to transport bump states
 50 on an $SO(2)$ -structured manifold; (ii) a Lie-algebraic transport analysis of the gain-imbalance
 51 mechanism; (iii) an analytic Fourier phase decoder that maps distributed ring activity to scalar phase
 52 estimates compatible with homeostatic comparison; (iv) a forward-kinematic Fourier decoder that
 53 reads out an $SO(3)$ orientation estimate from known revolute-chain ring states via product Fourier
 54 features; and (v) empirical validation in NEST spiking simulations [6] and on a Tiago robot [14]

55 The remainder of the paper presents the architecture, derives the $SO(2)$ transport mechanism and
 56 Fourier readouts, and evaluates the resulting spiking controller in NEST and on a Tiago arm. Addi-
 57 tional derivations and experimental details are provided in the supplementary material.

58 2 Architecture Overview

59 2.1 Single-ring architecture for revolute joint control




Figure 1: Single-ring architecture for revolute joint control, composed of four modules. The goal ring encodes the desired joint phase as a localized bump; the state ring encodes the current joint phase. A Positional Fourier readout decodes each bump into a scalar phase estimate, which feeds a Hammel-style homeostatic comparator. The comparator drives opponent left-gain (LG) and right-gain (RG) populations, moving the ring in their preferred direction. A mismatch between goal and state produces a gain imbalance that moves the bump state along the ring; when the mismatch vanishes, the opponent signals become balanced and stabilizes the bump.

This single-ring architecture builds on prior ring-attractor work for number subitization [8], extending the substrate from discretized representations of continuous revolute-joint control. Figure 1 shows the resulting control primitive for a single revolute joint, composed of goal and state ring-attractor populations, a Fourier readout layer, a homeostatic comparator, and opponent gain-modulation populations.

The goal ring encodes the desired joint configuration as the phase of a localized activity bump, and a similar representation is subsequently used by the state ring for the current joint configuration. The activity of these rings is projected onto positional Fourier features, which are decoded into spike-rate estimates of scalar phase. These goal and state phase estimates are then provided to a Hammel-inspired homeostatic comparator for feedback control [3]. The comparator produces opponent output signals that drive the left-gain (LG) and right-gain (RG) populations, directing the state bump along the preferred angular direction and closing the control loop.

This single-ring unit constitutes the basic control primitive. The following sections first derive the geometry and transport dynamics of this primitive, before Section 2.5 describes how multiple ring states are decoded through a known forward-kinematic chain.

2.2 Ring attractor as an $SO(2)$ -structured state space (Figure 1.A)

A ring attractor is a recurrent neural network whose connectivity is invariant to rotations of a circular preferred-phase index [23, 1, 5]. This symmetry supports a continuous family of stable population states: localized activity bumps centered at any phase $\phi \in S^1$. In the regime considered here, the controlled state is the bump centroid rather than the activity of any individual neuron.

For a revolute joint, this representation is natural because the joint angle and the bump phase both live on a circular manifold. We approximate the ring activity as a translated bump,

$$u_r(\theta, t) \approx f(\theta - \phi(t)), \quad (1)$$

where f is the bump profile and $\phi(t)$ is the represented joint phase. Thus, although the neural activity is high dimensional, the controlled coordinate is the scalar phase $\phi(t) \in S^1$.

This population code has an $SO(2)$ structure: shifting the bump by an angle α maps u_ϕ to $u_{\phi+\alpha}$. The family of bump states is therefore an orbit of the rotation group, and the centroid phase can equivalently be represented as $R(\phi) \in SO(2)$. Control corresponds to moving the bump along this orbit. The tangent direction is generated by the infinitesimal rotation operator $J = -\partial_\theta$, so an input that reliably changes the represented joint angle must project onto the bump tangent rather than merely perturbing firing rates.

In a symmetric ring, all centroid phases are equally stable and there is no preferred direction of drift. Balanced afferent drive preserves this symmetry and keeps the bump stationary. Directed motion requires breaking the symmetry through a signed imbalance between opponent directional drives, which produces a transport term of the form

$$\partial_t u_r(\theta, t) + c(u_r) \partial_\theta u_r(\theta, t) = F_r(u_r), \quad (2)$$

where $F_r(u_r)$ denotes the symmetric recurrent dynamics and $c(u_r)$ is a signed transport coefficient. The next section shows how opponent gain modulation implements this coefficient in the spiking circuit.

2.3 Gain modulation as Lie-algebraic transport (Figure 1.B)

Opponent gain modulation provides the directional drive that moves the bump. The architecture uses two gain populations, $u_g^+(\theta, t)$ and $u_g^-(\theta, t)$, which receive phase-shifted copies of the state-ring activity together with scalar homeostatic drives:

$$u_g^+(\theta, t) = W_{rg} g(u_r(\theta + \delta\theta, t)) + B_+(t), \quad u_g^-(\theta, t) = W_{rg} g(u_r(\theta - \delta\theta, t)) + B_-(t). \quad (3)$$

The opposite shifts make the two populations sensitive to opposite directions of bump motion. When $B_+(t) = B_-(t)$, their effects are balanced and the ring symmetry is preserved. When the homeostatic comparator creates an imbalance, the symmetry is broken and the bump drifts with a sign determined by $B_+(t) - B_-(t)$.

105 To first order in the phase shift $\delta\theta$, the summed gain feedback contains a symmetric component and a
 106 directional component proportional to the ring gradient:

$$g(u_g^+) + g(u_g^-) \approx \underbrace{g(x_+) + g(x_-)}_{\text{symmetric}} + \underbrace{W_{rg}\delta\theta g'(u_r)[g'(x_+) - g'(x_-)]\partial_\theta u_r}_{\text{directional}}, \quad (4)$$

107 where $x_\pm = W_{rg}g(u_r) + B_\pm(t)$. The symmetric term changes the net drive without selecting a
 108 direction. The directional term vanishes when the opponent drives are balanced and becomes nonzero
 109 only when $B_+(t) \neq B_-(t)$.

110 Collecting this directional feedback in the ring dynamics yields

$$\partial_t u_r(\theta, t) + c(u_r) \partial_\theta u_r(\theta, t) = F_r(u_r), \quad (5)$$

111 with transport coefficient

$$c(u_r) = \frac{W_{gr}W_{rg}\delta\theta}{\tau_r} g'(u_r)[g'(x_+) - g'(x_-)]. \quad (6)$$

112 Thus the transport speed scales with the phase shift, bidirectional coupling, and opponent gain
 113 difference. Since $J = -\partial_\theta$, this is equivalently activity injection along the tangent direction of the
 114 $SO(2)$ bump manifold:

$$\partial_t u_r(\theta, t) = F_r(u_r) + c(u_r) J u_r(\theta, t). \quad (7)$$

115 On the single-bump manifold, this transport reduces to an effective phase velocity,

$$\dot{\phi}(t) = c_{\text{eff}}(t),$$

116 where $c_{\text{eff}}(t)$ is the projection of the gain imbalance onto the bump tangent. Opponent gain modulation
 117 therefore encodes velocity along the ring, not joint angle itself; integrating this velocity transports the
 118 internal joint state toward the goal set by the homeostatic comparator.

119 2.4 Positional Fourier readout for scalar phase decoding (Figure 1.E)

120 The homeostatic comparator operates on scalar variables, whereas the ring attractor represents state
 121 as a distributed bump. We therefore decode each ring activity profile into a scalar phase estimate
 122 using a periodic Fourier readout.

123 For a translated bump $g(x) = r(x - \theta)$, the centroid estimator can be written as a projection of
 124 activity onto a positional ramp. Since the ring is periodic, this ramp admits a Fourier sine expansion,
 125 yielding the linear decoder

$$\hat{\theta} = b + \sum_{m=1}^K w_m S_m(g), \quad S_m(g) = \sum_{x=0}^N g(x) \sin\left(\frac{2\pi m x}{L}\right), \quad (8)$$

126 where w_m are analytic coefficients, optionally corrected for finite bump width. Appendix C gives the
 127 full derivation.

128 Figure 2 validates the decoder on translated Gaussian bumps and its spiking implementation. Increasing
 129 the number of modes improves phase reconstruction, while a single mode already captures the
 130 coarse phase signal. Because the features $S_m(g)$ are signed but spike counts are nonnegative, each
 131 feature is implemented with push-pull subpopulations and recovered by subtraction.

132 In the controller, the goal and state rings are decoded into $\hat{\theta}_g$ and $\hat{\theta}_s$. The Hammel comparator
 133 computes their signed mismatch and drives the opponent gain populations, thereby connecting the
 134 $SO(2)$ -structured population code to the gain-modulated transport mechanism.

135 2.5 Forward-kinematic Fourier decoder for known revolute chains

136 The single-ring controller represents and regulates one angular state. To connect multiple controlled
 137 rings to robot-level variables, we define a forward-kinematic readout for known revolute chains.
 138 This readout does not add a new control law; it maps distributed ring activity into an end-effector
 139 orientation estimate.




Figure 2: Analysis of the Fourier phase decoder. Top left: sine features $S_m(g)$ for Gaussian bump activity. Top right: Fourier decoders with $K = 1, 3, 50$ modes reconstruct centered position with increasing accuracy. Bottom left: push-pull spiking implementation of the $K = 1$ feature. Bottom right: mean $\pm$ standard deviation spike response across 10 repeats and 25 ring positions, showing a stable phase-dependent readout.

For each joint q_i , the Fourier readout provides first-harmonic estimates

$$\bar{C}_i \approx \cos q_i, \quad \bar{S}_i \approx \sin q_i,$$

which form the local feature vector

$$\varphi_i = \begin{bmatrix} 1 \\ \bar{C}_i \\ \bar{S}_i \end{bmatrix} \approx \begin{bmatrix} 1 \\ \cos q_i \\ \sin q_i \end{bmatrix}. \quad (9)$$

For a revolute joint with known axis, the corresponding rotation matrix is linear in these features. For a serial chain, products of joint rotations therefore generate products of local Fourier terms, all contained in

$$\Phi^{(n)} = \varphi_1 \otimes \varphi_2 \otimes \cdots \otimes \varphi_n. \quad (10)$$

Thus, for a fixed kinematic convention, there exists a linear map T such that

$$\text{vec}(\hat{R}_{ee}) = T\hat{\Phi}^{(n)}, \quad (11)$$

where $\hat{\Phi}^{(n)}$ is formed from ring-estimated features.

In the robot experiments, this decoder is used to read out lift, pitch, and yaw under a fixed ordered convention. In the spiking implementation, signed product-feature populations are formed from the input rings and projected to orientation-coordinate output rings, allowing the composition to remain within the neural substrate.

This decoder assumes known revolute kinematics, fixed rotation axes and multiplication order, and reliable first-mode sine-cosine estimates from each ring. Under these assumptions, it maps a product of local $SO(2)$ ring states to an $SO(3)$ orientation estimate, extending the single-joint control primitive toward task-space representations.




Figure 3: Robot-to-ring interface for the Tiago experiments. Each control step pairs a Tiago camera frame with the current spiking ring activity. The circular and unwrapped plots show the same population response, making the bump location and centroid visible during control. The timeline shows decoded gain-modulated commands moving the joint from start to target. Quantitative tracking results are reported in Figure 5, with additional analyses in Figures 7 and 8.

3 Implementation and Results

3.1 Experimental setup

We evaluate the architecture in two stages. Spiking experiments instantiate the full architecture in NEST [6] using `iaf_psc_alpha` populations and test whether the homeostatic gain-modulation loop transports a state bump toward a goal on the $SO(2)$ -structured manifold, with model parameters reported in Appendix Table 1. Robot experiments connect the resulting trajectories to a Tiago arm [14] via ROS, using a receding-horizon streaming interface to issue joint velocity commands and logging measured joint states to evaluate whether decoded ring phases correspond to physical motion. For multi-ring trials, end-effector orientation is computed from logged joint angles through the known forward-kinematic model.

3.2 Spiking implementation of $SO(2)$ homeostatic transport

We first evaluate the complete single-ring controller in NEST. The circuit contains the state ring attractor, the Hammel-style homeostatic comparator, and opponent left- and right-gain populations from Section 2. The goal is to verify the core mechanism from Section 2.3: a phase mismatch should induce an opponent gain imbalance, and this imbalance should transport the ring-attractor bump.

Figure 4 shows the resulting activity. The ring maintains a localized bump whose centroid is displaced from the start phase toward the goal phase, demonstrating directed transport of the represented state. The trajectory is not perfectly monotonic and does not always converge exactly to the target, but it remains a coherent bump motion rather than a diffuse change in population firing.

Panel C quantifies the remaining circular error across target locations after the transient period. Errors are target-dependent, indicating that the present parameterization produces a usable but nonuniform transport field over the ring. This nonuniformity is consistent with finite-size spiking effects, decoder resolution, and local gain calibration, and motivates the calibration refinements used in the robot interface.

Panel D shows the predicted opponent structure in the homeostatic and gain pathways. Hot and cold comparator activity varies systematically with relative phase, producing a graded mismatch signal.




Figure 4: Single-ring $SO(2)$ homeostatic controller implemented in NEST where the ring attractor is simulated using 200 neurons. **A**: ring-attractor firing-rate heatmap with the decoded circular centroid overlaid in white. The cyan dotted line marks the start phase and the orange dotted line marks the goal phase. **B**: left- and right-gain raster plots throughout the simulation period (associated to the top graph). **C**: relative goal error, computed as circular distance to the goal divided by ring size N , averaged over $t > 30000$ ms across ten runs per goal. The black dotted vertical line marks the starting point. **D**: left- and right-homeostasis neuron spike-count heatmaps.

181 The left- and right-gain channels show complementary activation, with each channel preferentially
 182 engaged for one sign of the phase mismatch. This push-pull pattern implements the imbalance
 183 $B_+(t) - B_-(t)$ that generates the transport coefficient in Equation (12). Together, these results
 184 show that the spiking circuit realizes the intended symmetry-breaking mechanism for $SO(2)$ bump
 185 transport, while also revealing the target-dependent residual error expected from an unoptimized
 186 finite spiking implementation.

187 3.3 Tiago single-joint control from gain-modulated ring velocity

188 We next evaluate whether the $SO(2)$ ring controller can drive physical joint motion on the Tiago arm.
 189 The central test is whether opponent gain modulation produces a signed velocity command whose
 190 integration moves a revolute joint toward a target angle.

191 For each Tiago arm joint, the measured joint angle is mapped onto the state ring and a sampled
 192 target angle is injected into the goal ring. The NEST controller runs in closed loop; at each step, the
 193 difference in spike counts between the right-gain and left-gain populations,

$$s_t = R_t - L_t, \quad (12)$$

194 provides the signed gain signal. This is the spiking realisation of the gain imbalance derived in
 195 Section 2.3: positive s_t drives transport in one angular direction, negative s_t in the other. The signal
 196 is passed through an exponential filter and decoded into a joint velocity command with separate
 197 calibration gains for the positive and negative drive components, preserving the opponent structure
 198 of the controller. Commands are sent as receding-horizon joint trajectories via ROS, and measured
 199 joint positions are logged throughout. Figure 5 summarises the final absolute error across all seven
 200 Tiago arm joints. The controller successfully drives directed motion across all joints, confirming that
 201 the gain-modulated ring signal translates to physical displacement. Performance is joint-dependent:




Figure 5: Tiago single-joint control where the opponent gain-modulation activity is translated to the robot’s joint velocity. Left: Robot simulation snapshot demonstrating how the robot moves for specific joint index. Right: Robot simulation used to demonstrate error between goal state and target state of the ring attractor based single joint control.

some joints show consistently small errors across batches, while others exhibit larger spread and occasional outliers. This variability reflects differences in joint range, mechanical dynamics, and the calibration of the neural drive to physical velocity, rather than failures of the control mechanism itself.

3.4 NEST implementation of the forward-kinematic ring decoder

We finally evaluate whether the product-feature readout derived in Section 2.5 can be implemented in spiking neurons and decoded into orientation-coordinate output rings. We instantiate a two-ring decoder for a known two-revolute-joint chain,

$$R(q_1, q_2) = R_x(q_1) R_y(q_2), \quad (13)$$

where the first ring represents rotation about the x -axis and the second about the y -axis. Although only two input rings are used, the experiment tests the essential compositional mechanism: local $SO(2)$ Fourier features are combined into signed product-feature populations via push-pull pairs, and those product features are projected to output rings encoding lift, pitch, and yaw.

Figure 10 shows the internal decoder activity. During a sweep of q_2 with q_1 fixed, the first-order features track the sweep directly while the product features reflect the interaction between the fixed q_1 state and the varying q_2 state. The output ring profiles confirm that these product features are sufficient to produce structured orientation-coordinate population responses.

Figure 6 quantifies the decoder output against analytically computed targets across two configuration-space slices. The decoded traces closely track the true targets throughout, and Figure 9 shows that the circular index error remains small across both sweep conditions and all three orientation variables. The discontinuities visible in Figure 6 are expected sawtooth wrap-around effects, not decoder failures.

These results validate the orientation readout in spiking neurons. The decoder does not control orientation — it reads out the orientation implied by the ring states. The single-ring homeostatic gain-modulation loop generates $SO(2)$ trajectories; the multi-ring product-feature decoder composes those trajectories into lift, pitch, and yaw output rings.




Figure 6: Sawtooth-decoded output-ring index against the true target index for two slices of the joint configuration space (q_1 fixed with q_2 varied; q_2 fixed with q_1 varied). Decoded and true traces closely overlap for lift, pitch, and yaw. Discontinuities are expected sawtooth wrap-around effects from the circular representation, quantified in Figure 9. See Appendix Figure 10 for decoder activity.

227 4 Conclusion

228 We introduced a spiking ring-attractor architecture for representing and controlling revolute robot
 229 joints. Rather than treating joint angle as an external scalar variable passed to a conventional controller,
 230 the proposed system represents joint state as a persistent activity bump on a circular neural manifold
 231 and regulates that state through homeostatic comparison and opponent gain modulation. This provides
 232 a different paradigm for spiking robot control: representation and action are not separated, but are
 233 coupled through transport of a population code whose geometry matches the topology of the physical
 234 joint.

235 The results show that gain imbalance can induce directed bump motion on an $SO(2)$ -structured ring,
 236 that distributed ring activity can be decoded through analytic Fourier features, and that local ring
 237 features can be composed into forward-kinematic orientation readouts for known revolute chains.
 238 Robot experiments further demonstrate that the gain-modulated ring dynamics can be coupled to
 239 physical joint motion, producing directed actuation with joint-dependent accuracy. Although the
 240 present implementation is not intended to replace optimized model-based or PID controllers, it
 241 establishes a neuromorphic control primitive in which circular state representation, regulation, and
 242 motor drive are implemented within a common spiking substrate.

243 Because the controlled variable is represented by a population activity pattern rather than a single
 244 scalar state, this framework may offer robustness advantages under noise, partial perturbation, or
 245 neuronal variability: errors in individual units need not destroy the represented phase as long as
 246 the population bump remains coherent. More broadly, precise control of such attractor bumps
 247 could enable applications beyond low-level joint actuation, including task-space control, internal
 248 simulation, and neural forms of motor imagery. In biological systems, structured activity patterns
 249 can be maintained, shifted, and replayed offline, as in hippocampal experience replay; analogous
 250 mechanisms in robotic ring-attractor substrates could support imagined trajectories, predictive rollouts,
 251 or planning over embodied state spaces. We therefore view controlled spiking attractor dynamics as a
 252 promising route toward natural neural controllers for revolute systems, where the geometry of the
 253 controller is aligned with the geometry of the body it controls.

References

- [1] Shun-ichi Amari. Dynamics of pattern formation in lateral-inhibition type neural fields. *Biological Cybernetics*, 27(2):77–87, 1977. ISSN 1432-0770. doi: 10.1007/bf00337259. URL <http://dx.doi.org/10.1007/BF00337259>.
- [2] Sabrina Benas, Ximena Fernandez, and Emilio Kropff. Modeled grid cells aligned by a flexible attractor. May 2024. doi: 10.7554/elife.89851.2. URL <http://dx.doi.org/10.7554/elife.89851.2>.
- [3] Jack A. Boulant. Neuronal basis of hammel’s model for set-point thermoregulation. *Journal of Applied Physiology*, 100(4):1347–1354, 2006. doi: 10.1152/japplphysiol.01064.2005.
- [4] Chris Eliasmith. A unified approach to building and controlling spiking attractor networks. *Neural Computation*, 17(6):1276–1314, 2005. doi: 10.1162/0899766053630332.
- [5] Wulfram Gerstner, Werner M. Kistler, Richard Naud, and Liam Paninski. *Cortical field models for perception*, page 467–490. Cambridge University Press, 2014.
- [6] Marc-Oliver Gewaltig and Markus Diesmann. Nest (neural simulation tool). *Scholarpedia*, 2(4):1430, 2007. doi: 10.4249/scholarpedia.1430.
- [7] Jonathan Green, Atsuko Adachi, Kunal K. Shah, Jonathan D. Hirokawa, Pablo S. Magani, and Gaby Maimon. A neural circuit architecture for angular integration in drosophila. *Nature*, 546(7656):101–106, May 2017. ISSN 1476-4687. doi: 10.1038/nature22343. URL <http://dx.doi.org/10.1038/nature22343>.
- [8] Aung Htet, Alejandro Jimenez Rodriguez, Sarah Hamburg, and Alessandro Di Nuovo. *Allostatic Control of Persistent States in Spiking Neural Networks for Perception and Computation*, page 388–402. Springer Nature Switzerland, November 2025. ISBN 9783032074485. doi: 10.1007/978-3-032-07448-5_33. URL http://dx.doi.org/10.1007/978-3-032-07448-5_33.
- [9] Mikail Khona and Ila R. Fiete. Attractor and integrator networks in the brain. *Nature Reviews Neuroscience*, 23(12):744–766, November 2022. ISSN 1471-0048. doi: 10.1038/s41583-022-00642-0. URL <http://dx.doi.org/10.1038/s41583-022-00642-0>.
- [10] Sung Soo Kim, Hervé Rouault, Shaul Druckmann, and Vivek Jayaraman. Ring attractor dynamics in the *drosophila* central brain. *Science*, 356(6340):849–853, May 2017. ISSN 1095-9203. doi: 10.1126/science.aal4835. URL <http://dx.doi.org/10.1126/science.aal4835>.
- [11] Thomas C. Knowles, Anna G. Summerton, James G. H. Whiting, and Martin J. Pearson. Ring attractors as the basis of a biomimetic navigation system. *Biomimetics*, 8(5), 2023. ISSN 2313-7673. doi: 10.3390/biomimetics8050399. URL <https://www.mdpi.com/2313-7673/8/5/399>.
- [12] Richard M. Murray, Zexiang Li, and S. Shankar Sastry. *Rigid Body Motion*, page 19–79. CRC Press, December 2017. ISBN 9781315136370. doi: 10.1201/9781315136370-2. URL <http://dx.doi.org/10.1201/9781315136370-2>.
- [13] Peter Mussells Pires, Lingwei Zhang, Victoria Parache, L. F. Abbott, and Gaby Maimon. Converting an allocentric goal into an egocentric steering signal. *Nature*, 626(8000):808–818, February 2024. ISSN 1476-4687. doi: 10.1038/s41586-023-07006-3. URL <http://dx.doi.org/10.1038/s41586-023-07006-3>.
- [14] Jordi Pagès, Luca Marchionni, and Francesco Ferro. Tiago: the modular robot that adapts to different research needs. 2016. URL <https://api.semanticscholar.org/CorpusID:218478582>.
- [15] Jesús D. Rivero-Ortega, Juan S. Mosquera-Maturana, Josh Pardo-Cabrera, Julián Hurtado-López, Juan D. Hernández, Víctor Romero-Cano, and David F. Ramírez-Moreno. Ring attractor bio-inspired neural network for social robot navigation. *Frontiers in Neurobotics*, Volume 17 - 2023, 2023. ISSN 1662-5218. doi: 10.3389/fnbot.2023.1211570. URL <https://www.frontiersin.org/journals/neurobotics/articles/10.3389/fnbot.2023.1211570>.

- [16] Emilio Salinas and Terrence J. Sejnowski. Gain modulation in the central nervous system: Where behavior, neurophysiology, and computation meet. *The Neuroscientist*, 7(5):430–440, October 2001. ISSN 1089-4098. doi: 10.1177/107385840100700512. URL <http://dx.doi.org/10.1177/107385840100700512>.
- [17] Johannes D. Seelig and Vivek Jayaraman. Neural dynamics for landmark orientation and angular path integration. *Nature*, 521(7551):186–191, May 2015. ISSN 1476-4687. doi: 10.1038/nature14446. URL <http://dx.doi.org/10.1038/nature14446>.
- [18] Daniel Turner-Evans, Stephanie Wegener, Hervé Rouault, Romain Franconville, Tanya Wolff, Johannes D Seelig, Shaul Druckmann, and Vivek Jayaraman. Angular velocity integration in a fly heading circuit. *eLife*, 6, May 2017. ISSN 2050-084X. doi: 10.7554/elife.23496. URL <http://dx.doi.org/10.7554/eLife.23496>.
- [19] Pantelis Vafidis, David Oswald, Tiziano D’Albis, and Richard Kempter. Learning accurate path integration in ring attractor models of the head direction system. *eLife*, 11:e69841, jun 2022. ISSN 2050-084X. doi: 10.7554/eLife.69841. URL <https://doi.org/10.7554/eLife.69841>.
- [20] Elena A. Westeinde, Emily Kellogg, Paul M. Dawson, Jenny Lu, Lydia Hamburg, Benjamin Midler, Shaul Druckmann, and Rachel I. Wilson. Transforming a head direction signal into a goal-oriented steering command. *Nature*, 626(8000):819–826, February 2024. ISSN 1476-4687. doi: 10.1038/s41586-024-07039-2. URL <http://dx.doi.org/10.1038/s41586-024-07039-2>.
- [21] Si Wu, Ying Wu, Wen-Hao Zhang, and Junfeng Zuo. The motion planning neural circuit in goal-directed navigation as lie group operator search. In *Advances in Neural Information Processing Systems 37*, NeurIPS 2024, page 110369–110387. Neural Information Processing Systems Foundation, Inc. (NeurIPS), 2024. doi: 10.52202/079017-3503. URL <http://dx.doi.org/10.52202/079017-3503>.
- [22] Jian Zhang and L.F. Abbott. Gain modulation of recurrent networks. *Neurocomputing*, 32-33:623–628, 2000. ISSN 0925-2312. doi: 10.1016/S0925-2312(00)00224-1. URL [http://dx.doi.org/10.1016/S0925-2312\(00\)00224-1](http://dx.doi.org/10.1016/S0925-2312(00)00224-1).
- [23] K Zhang. Representation of spatial orientation by the intrinsic dynamics of the head-direction cell ensemble: a theory. *The Journal of Neuroscience*, 16(6):2112–2126, March 1996. ISSN 1529-2401. doi: 10.1523/jneurosci.16-06-02112.1996. URL <http://dx.doi.org/10.1523/JNEUROSCI.16-06-02112.1996>.

A Ring-attractor mode analysis and SO(2) reduction

This appendix provides the derivation underlying Section 2.2. We start from the continuum ring neural field

$$\tau_r \partial_t u_r(\theta, t) = -u_r(\theta, t) + \int_{-\pi}^{\pi} W(\theta - \theta') g(u_r(\theta', t)) d\theta' + I(\theta, t),$$

where $u_r(\theta, t)$ is the ring activity, $W(\theta - \theta')$ is a translation-invariant recurrent kernel, g is a static nonlinearity, and I is external input. For the mode analysis, we first consider a constant input $I(\theta, t) = I_0$ and a spatially homogeneous equilibrium

$$u_r(\theta, t) = \bar{u}.$$

Substituting this into the neural field and setting $\partial_t u_r = 0$ gives

$$\bar{u} = g(\bar{u}) \int_{-\pi}^{\pi} W(\theta - \theta') d\theta' + I_0.$$

Since the kernel depends only on the phase difference, the integral is independent of θ . Let

$$\widehat{W}_0 = \int_{-\pi}^{\pi} W(\psi) d\psi.$$

Then the homogeneous fixed point satisfies

$$\bar{u} = g(\bar{u}) \widehat{W}_0 + I_0.$$

To study the stability of this homogeneous state, write

$$u_r(\theta, t) = \bar{u} + v(\theta, t),$$

where v is a small perturbation. Expanding the nonlinearity to first order,

$$g(\bar{u} + v) \approx g(\bar{u}) + g'(\bar{u})v.$$

Substituting into the neural field and cancelling the fixed-point terms gives the linearized dynamics

$$\tau_r \partial_t v(\theta, t) = -v(\theta, t) + g'(\bar{u}) \int_{-\pi}^{\pi} W(\theta - \theta') v(\theta', t) d\theta'.$$

Because the recurrent operator is convolutional on the ring, Fourier modes diagonalize the linearized dynamics. Let

$$v(\theta, t) = \sum_{k=-\infty}^{\infty} a_k(t) e^{ik\theta}.$$

The Fourier coefficient of the kernel is

$$\widehat{W}_k = \int_{-\pi}^{\pi} W(\psi) e^{-ik\psi} d\psi.$$

Using the convolution identity on the ring, each mode evolves independently:

$$\tau_r \dot{a}_k(t) = \left[-1 + g'(\bar{u}) \widehat{W}_k \right] a_k(t).$$

Thus the growth rate of mode k is

$$\lambda_k = \frac{1}{\tau_r} \left[-1 + g'(\bar{u}) \widehat{W}_k \right].$$

Modes with $\lambda_k < 0$ decay, while modes with $\lambda_k > 0$ destabilize the homogeneous state. In the parameter regime considered in the main text, the recurrent kernel and nonlinearity support a stable single-bump profile rather than a homogeneous firing state. The mode analysis therefore motivates the reduction from a full activity field to a low-dimensional bump coordinate.

We assume the ring is operating in this activated single-bump regime. The activity can then be approximated by a fixed bump profile translated by a phase coordinate,

$$u_r(\theta, t) \approx f(\theta - \phi(t)),$$

where f is the bump profile and $\phi(t) \in S^1$ is the bump centroid. This is a collective-coordinate approximation: the controlled state is not the activity of a single neuron, but the position of the distributed bump.

The rotation action on ring activity is

$$(\rho_\alpha u_r)(\theta) = u_r(\theta - \alpha).$$

For the bump state

$$u_\phi(\theta) = f(\theta - \phi),$$

we have

$$(\rho_\alpha u_\phi)(\theta) = f(\theta - \alpha - \phi) = u_{\phi+\alpha}(\theta).$$

Thus the family of translated bump states is an orbit of the continuous rotation group. The bump centroid may equivalently be represented by the rotation matrix

$$R(\phi) = \begin{pmatrix} \cos \phi & -\sin \phi \\ \sin \phi & \cos \phi \end{pmatrix} \in SO(2).$$

The full neural population state lives in a high-dimensional activity space, but the controlled collective coordinate $\phi(t)$ lies on S^1 and carries the action of $SO(2)$.

The infinitesimal generator of rotations on the ring is

$$J = -\partial_\theta.$$

For a small displacement $\delta\phi$,

$$u_{\phi+\delta\phi}(\theta) = f(\theta - \phi - \delta\phi) \approx f(\theta - \phi) - \delta\phi \partial_\theta f(\theta - \phi).$$

Using $J = -\partial_\theta$, this becomes

$$u_{\phi+\delta\phi}(\theta) \approx u_\phi(\theta) + \delta\phi J u_\phi(\theta).$$

Thus the tangent direction of the single-bump manifold is $J u_\phi$. A control input that moves the bump should therefore project onto this direction.

B Derivation of gain modulation as a transport equation

This appendix provides the derivation underlying Section 2.3. The goal is to show how opponent gain populations produce a first-order transport term on the ring. Let $u_g^+(\theta, t)$ and $u_g^-(\theta, t)$ denote the two gain populations. We model their inputs as shifted copies of the ring activity plus opponent homeostatic drive:

$$u_g^+(\theta, t) = W_{rg} g(u_r(\theta + \delta\theta, t)) + B_+(t),$$

$$u_g^-(\theta, t) = W_{rg} g(u_r(\theta - \delta\theta, t)) + B_-(t),$$

where W_{rg} is the ring-to-gain connection strength, $\delta\theta$ is a small connection phase shift, and $B_+(t)$, $B_-(t)$ are homeostatic drive terms.

For small $\delta\theta$, expand the shifted ring activity:

$$u_r(\theta + \delta\theta, t) \approx u_r(\theta, t) + \delta\theta \partial_\theta u_r(\theta, t),$$

$$u_r(\theta - \delta\theta, t) \approx u_r(\theta, t) - \delta\theta \partial_\theta u_r(\theta, t).$$

Applying the nonlinearity gives

$$g(u_r(\theta + \delta\theta, t)) \approx g(u_r) + \delta\theta g'(u_r) \partial_\theta u_r,$$

$$g(u_r(\theta - \delta\theta, t)) \approx g(u_r) - \delta\theta g'(u_r) \partial_\theta u_r,$$

where u_r abbreviates $u_r(\theta, t)$.

387 Define the baseline gain inputs

$$x_+ = W_{rg}g(u_r) + B_+(t), \quad x_- = W_{rg}g(u_r) + B_-(t).$$

388 Then the gain-population activities can be written to first order as

$$u_g^+ \approx x_+ + W_{rg}\delta\theta g'(u_r)\partial_\theta u_r,$$

389

$$u_g^- \approx x_- - W_{rg}\delta\theta g'(u_r)\partial_\theta u_r.$$

390 Expanding the gain output nonlinearity around x_+ and x_- gives

$$g(u_g^+) \approx g(x_+) + W_{rg}\delta\theta g'(x_+)g'(u_r)\partial_\theta u_r,$$

391

$$g(u_g^-) \approx g(x_-) - W_{rg}\delta\theta g'(x_-)g'(u_r)\partial_\theta u_r.$$

392 Summing the two opponent outputs gives

$$g(u_g^+) + g(u_g^-) \approx g(x_+) + g(x_-) + W_{rg}\delta\theta g'(u_r) [g'(x_+) - g'(x_-)] \partial_\theta u_r.$$

393 The first two terms are nondirectional. They modify the local drive to the ring but do not select a
394 direction of motion. The final term is directional because it is proportional to $\partial_\theta u_r$.

395 This expression also shows why balanced opponent drive does not move the bump. If $B_+(t) = B_-(t)$,
396 then $x_+ = x_-$, and therefore

$$g'(x_+) - g'(x_-) = 0.$$

397 The directional term vanishes. Thus equal opponent magnitudes preserve the rotational symmetry
398 of the ring and do not generate bump transport. Directed motion requires $B_+(t) \neq B_-(t)$, which
399 breaks the symmetry and leaves a nonzero first-order transport term.

400 The state ring receives feedback from the gain populations with strength W_{gr} . Substituting the
401 directional term into the ring equation and collecting all nondirectional terms into $F_r(u_r)$ gives

$$\partial_t u_r(\theta, t) + c(u_r)\partial_\theta u_r(\theta, t) = F_r(u_r).$$

402 The symmetric part is

$$F_r(u_r) = \frac{1}{\tau_r} \left[-u_r(\theta, t) + \int_{-\pi}^{\pi} W(\theta - \theta')g(u_r(\theta', t)) d\theta' + I(\theta, t) \right].$$

403 The transport coefficient is

$$c(u_r) = \frac{W_{gr}W_{rg}\delta\theta}{\tau_r} g'(u_r) [g'(x_+) - g'(x_-)].$$

404 Depending on the sign convention used for the feedback projection, the sign of $c(u_r)$ may be absorbed
405 into W_{gr} or into the definitions of the positive and negative gain channels. The important point is that
406 the transport coefficient is proportional to the phase shift $\delta\theta$, the bidirectional coupling $W_{gr}W_{rg}$, and
407 the local gain difference between the two opponent channels.

408 Using $J = -\partial_\theta$, the transport equation becomes

$$\partial_t u_r(\theta, t) = F_r(u_r) + c(u_r)Ju_r(\theta, t).$$

409 Thus gain modulation supplies the scalar coefficient multiplying the Lie algebra generator of the ring.

410 Finally, consider the single-bump approximation

$$u_r(\theta, t) = f(\theta - \phi(t)).$$

411 Then

$$\partial_t u_r(\theta, t) = -\dot{\phi}(t)f'(\theta - \phi(t)) = \dot{\phi}(t)Ju_r(\theta, t).$$

412 If $c(u_r)$ is approximately constant over the bump support, or if it is projected onto the bump tangent
413 direction, the dynamics reduce to an effective phase equation

$$\dot{\phi}(t) = c_{\text{eff}}(t),$$

414 where $c_{\text{eff}}(t)$ is the transport coefficient averaged or projected over the bump. Therefore opponent
415 gain modulation encodes velocity along the $SO(2)$ -structured bump manifold, rather than encoding
416 the joint angle directly.

417 C Fourier decoder derivation

418 This appendix provides the derivation underlying the positional Fourier decoder in Section 2.4. The
 419 problem is to convert a distributed bump state into a scalar phase estimate that can be used by the
 420 homeostatic comparator.

421 Let the ring activity be a translated bump profile

$$g(x) = r(x - \theta),$$

422 where $x = 0, \dots, N$ indexes discrete ring positions and θ is the bump position. The centroid estimator
 423 is

$$\hat{\theta}_{\text{cent}} = \frac{\sum_{x=0}^N xg(x)}{\sum_{x=0}^N g(x)}.$$

424 Assume that the bump shape and amplitude are approximately constant under translation. Then the
 425 total mass

$$M_0 = \sum_{x=0}^N g(x)$$

426 is approximately independent of θ . The centroid can then be written as the linear decoder

$$\hat{\theta} \approx \frac{1}{M_0} \sum_{x=0}^N xg(x).$$

427 To center the decoded coordinate, define

$$L = N + 1, \quad \xi(x) = x - \frac{L}{2}.$$

428 Then

$$\hat{\theta} - \frac{L}{2} \approx \frac{1}{M_0} \sum_{x=0}^N \xi(x)g(x).$$

429 Thus position decoding is equivalent to projecting the activity profile onto a centered ramp function.

430 On a periodic interval, the centered ramp is represented by a sine series. For $x \in [0, L]$,

$$x - \frac{L}{2} = -\frac{L}{\pi} \sum_{m=1}^{\infty} \frac{1}{m} \sin\left(\frac{2\pi mx}{L}\right).$$

431 Substituting this into the centered centroid estimator gives

$$\hat{\theta} - \frac{L}{2} \approx \frac{1}{M_0} \sum_{x=0}^N g(x) \left[-\frac{L}{\pi} \sum_{m=1}^{\infty} \frac{1}{m} \sin\left(\frac{2\pi mx}{L}\right) \right].$$

432 Interchanging the order of summation yields

$$\hat{\theta} - \frac{L}{2} \approx -\frac{L}{\pi M_0} \sum_{m=1}^{\infty} \frac{1}{m} \sum_{x=0}^N g(x) \sin\left(\frac{2\pi mx}{L}\right).$$

433 Define the m -th sine feature

$$S_m(g) = \sum_{x=0}^N g(x) \sin\left(\frac{2\pi mx}{L}\right).$$

434 Then

$$\hat{\theta} \approx \frac{L}{2} - \frac{L}{\pi M_0} \sum_{m=1}^{\infty} \frac{1}{m} S_m(g).$$

435 Truncating the series to K modes gives the practical decoder

$$\hat{\theta} = b + \sum_{m=1}^K w_m S_m(g),$$

436 with

$$b = \frac{L}{2}, \quad w_m = -\frac{L}{\pi M_0 m}.$$

437 If the bump profile has finite width, then each sine feature is attenuated by the Fourier response of the
438 bump shape. In this case,

$$S_m(g) \approx A_m \sin\left(\frac{2\pi m \theta}{L}\right),$$

439 where A_m depends on the shape and width of the bump. A shape-corrected decoder can therefore
440 absorb this attenuation into the weights:

$$w_m = -\frac{L}{\pi m A_m}.$$

441 This shows that the Fourier decoder is a periodic reformulation of the centroid estimator. The centroid
442 projects activity onto a ramp; the Fourier decoder expresses that ramp in a sine basis adapted to the
443 ring geometry.

444 **D Signed implementation of Fourier features**

445 The sine features used by the decoder can take positive or negative values, whereas firing rates and
446 spike counts are nonnegative. To implement signed Fourier features with nonnegative neural activity,
447 each sine feature is represented by two subpopulations:

$$S_m^+(x) = \max\left(\sin\left(\frac{2\pi m x}{L}\right), 0\right),$$

448

$$S_m^-(x) = \max\left(-\sin\left(\frac{2\pi m x}{L}\right), 0\right).$$

449 The signed feature is recovered by subtraction:

$$S_m(g) = S_m^+(g) - S_m^-(g).$$

450 In practice, the two subpopulations accumulate spike responses to the positive and negative parts
451 of the sine feature. The downstream linear decoder then combines the recovered signed features
452 to estimate the bump phase. This implements the Fourier phase readout using nonnegative neural
453 responses while preserving the linear decoding structure used in the main text.

454 **E Compositional $SO(2)$ ring decoder for $SO(3)$ orientation**

455 This appendix derives the orientation readout used in the robot experiments. The control mechanism
456 — homeostatic comparison driving opponent gain modulation on individual rings — is described in
457 the main text. Here we show how the decoded phases of several such rings compose into a robot-level
458 $SO(3)$ orientation estimate through forward kinematics. This is a readout, not an additional control
459 law.

460 **E.1 From ring activity to local trigonometric features**

461 The Fourier readout derived in Appendix C extracts a scalar phase estimate from each ring. For the
462 orientation readout we retain the trigonometric form rather than collapsing to a scalar angle, because
463 forward kinematics is naturally expressed in terms of sines and cosines. For ring i with activity profile
464 $g_i(x)$ over positions $x = 0, \dots, L-1$, the first harmonic Fourier features are

$$C_i = \sum_{x=0}^{L-1} g_i(x) \cos\left(\frac{2\pi x}{L}\right), \quad S_i = \sum_{x=0}^{L-1} g_i(x) \sin\left(\frac{2\pi x}{L}\right). \quad (14)$$

465 When bump shape and amplitude are approximately preserved under translation — the same condition
466 that justifies the centroid approximation in the main text — normalising by the bump shape gives
467 reliable estimates

$$\bar{C}_i \approx \cos q_i, \quad \bar{S}_i \approx \sin q_i, \quad (15)$$

where q_i is the phase represented by the bump centroid. These are collected into a local feature vector

$$\varphi_i = \begin{bmatrix} 1 \\ \bar{C}_i \\ \bar{S}_i \end{bmatrix} \approx \begin{bmatrix} 1 \\ \cos q_i \\ \sin q_i \end{bmatrix}, \quad (16)$$

where the constant term accounts for the fixed entries present in all coordinate-axis rotation matrices.

E.2 A single revolute rotation is linear in ring features

A rotation about a coordinate axis $a_i \in \{x, y, z\}$ takes the form

$$R_x(q_i) = \begin{bmatrix} 1 & 0 & 0 \\ 0 & \cos q_i & -\sin q_i \\ 0 & \sin q_i & \cos q_i \end{bmatrix}, \quad R_y(q_i) = \begin{bmatrix} \cos q_i & 0 & \sin q_i \\ 0 & 1 & 0 \\ -\sin q_i & 0 & \cos q_i \end{bmatrix}, \quad R_z(q_i) = \begin{bmatrix} \cos q_i & -\sin q_i & 0 \\ \sin q_i & \cos q_i & 0 \\ 0 & 0 & 1 \end{bmatrix}. \quad (17)$$

Each matrix is linear in $1, \cos q_i, \sin q_i$, so for any coordinate axis the local rotation can be written as

$$R_{a_i}(q_i) = B_{a_i}^{(0)} + \cos q_i B_{a_i}^{(c)} + \sin q_i B_{a_i}^{(s)}, \quad (18)$$

where the coefficient matrices $B^{(0)}, B^{(c)}, B^{(s)}$ are read off directly from equation (17) and are fixed once the joint axis is known. A single ring's Fourier features therefore suffice to linearly represent one coordinate-axis revolute rotation.

E.3 Composing rings produces product Fourier features

For a serial chain of n revolute joints, the end-effector orientation is the ordered product

$$R_{ee}(q_1, \dots, q_n) = R_{a_1}(q_1) R_{a_2}(q_2) \cdots R_{a_n}(q_n). \quad (19)$$

The order here is not arbitrary — serial rotations do not commute, so the multiplication order is part of the kinematic convention and must match the robot model. Substituting equation (18) for each factor and expanding yields matrix entries that are sums of products of local sine and cosine terms across joints, such as

$$\cos q_1 \cos q_2, \quad \cos q_1 \sin q_2, \quad \sin q_1 \sin q_2, \quad \cos q_1 \cos q_2 \cos q_3. \quad (20)$$

All such terms appear in the tensor product of the individual feature vectors. The chain feature vector

$$\Phi^{(n)} = \varphi_1 \otimes \varphi_2 \otimes \cdots \otimes \varphi_n \quad (21)$$

contains every cross-term generated by the kinematic chain and has 3^n entries for n first-harmonic rings. Since every entry of R_{ee} is a sum of these product features, there exists a fixed matrix T such that

$$\text{vec}(R_{ee}) = T\Phi^{(n)}, \quad (22)$$

which is exact for coordinate-axis revolute chains. For the three-ring case considered here $\Phi^{(n)}$ has $3^3 = 27$ entries, which is entirely manageable; longer chains would require feature pruning or low-rank approximation.

E.4 Neural implementation with estimated features

In the spiking model the exact features are replaced by ring-activity estimates, giving

$$\hat{\Phi}^{(n)} = \hat{\varphi}_1 \otimes \hat{\varphi}_2 \otimes \cdots \otimes \hat{\varphi}_n, \quad \text{vec}(\hat{R}_{ee}) = T\hat{\Phi}^{(n)}. \quad (23)$$

Because product Fourier features can take positive or negative values while neural firing rates are nonnegative, each feature is implemented via a push-pull pair of subpopulations,

$$\Phi_j^+ = \max(\Phi_j, 0), \quad \Phi_j^- = \max(-\Phi_j, 0), \quad \Phi_j = \Phi_j^+ - \Phi_j^-, \quad (24)$$

preserving the linear readout structure while remaining compatible with nonnegative spike counts. The signed product populations project directly to output rings for lift, pitch, and yaw using analytic or learned weights.

496 E.5 Three-ring orientation readout for lift, pitch, and yaw

497 For the robot experiments three ring coordinates represent lift, pitch, and yaw. Under the ordered
498 convention

$$R = R_{\text{lift}}(\hat{q}_{\text{lift}}) R_{\text{pitch}}(\hat{q}_{\text{pitch}}) R_{\text{yaw}}(\hat{q}_{\text{yaw}}), \quad (25)$$

499 the composed feature vector $\Phi = \varphi_{\text{lift}} \otimes \varphi_{\text{pitch}} \otimes \varphi_{\text{yaw}}$ supports the exact linear readout $\text{vec}(R) = T\Phi$.
500 When an explicit rotation matrix is decoded, lift, pitch, and yaw are recovered as

$$\hat{\theta} = -\arcsin(r_{31}), \quad \hat{\lambda} = \arctan 2(r_{32}, r_{33}), \quad \hat{\psi} = \arctan 2(r_{21}, r_{11}), \quad (26)$$

501 valid away from the gimbal singularity at $\cos \theta = 0$, which does not arise in the workspace used for
502 the Tiago experiments. Alternatively, a direct linear decoder from $\hat{\Phi}$ to each output ring coordinate
503 can be fitted, bypassing explicit matrix extraction and allowing the composition to be carried out
504 entirely within the neural substrate.

505 The $SO(3)$ orientation produced by this readout is not separately controlled. It is the forward-
506 kinematic image of three independently controlled $SO(2)$ ring trajectories, each driven by its own
507 homeostatic gain-modulation loop.

508 F Extra Results

509 F.1 Tiago single-joint control from gain-modulated ring velocity




Figure 7: Left: Gain-population imbalance versus decoded velocity command across joints 0,3 and 5. Each point represents one recorded time step from the ring controller, showing an approximately linear increase in decoded velocity with the difference between right and left gain-population spike counts. Right: Settling time for joints 0, 3, and 5 under the ring and PID controllers. Settling time is defined as the first time the absolute tracking error remains below 0.05 rad for 5 consecutive samples; bars show mean $\pm$ 1 SD across trials.




Figure 8: Left: Absolute tracking error over time for joints 0, 3, and 5, shown for representative Ring and PID trials. The dashed line indicates the 0.05 rad tolerance. PID drops into this band quickly, while the Ring controller reduces error more slowly and reaches the same threshold later. Right: Velocity command over time for joints 0, 3, and 5, shown for representative Ring and PID trials. PID starts with a strong corrective command and rapidly decays toward zero, while the Ring controller applies smaller, smoother commands that vary gradually over a longer duration.




Figure 9: Circular index error for the two-ring forward-kinematic decoder under the same sweep conditions as Figure 6. Across both sweeps and all three decoded orientation variables, the circular error remains small, confirming that the discontinuities in Figure 6 arise from the circular representation rather than decoder failure.




Figure 10: Spiking implementation of the two-ring forward-kinematic decoder in NEST. Top: signed product-feature population activity during a sweep of q_2 with q_1 fixed, showing systematic variation of first-order features ($\cos q_1$, $\sin q_1$, $\cos q_2$, $\sin q_2$) and second-order products ($\cos q_1 \cos q_2$, etc.) with joint angle. Bottom: output ring profiles for lift, pitch, and yaw at representative q_2 values, showing structured population responses driven by the product-feature layer.

Table 1: Spiking model parameters used for the ring, homeostatic comparator, and gain-modulation populations.

Parameter	Symbol	Ring	Homeostatic	Gain
<i>Neuron parameters for iaf_psc_alpha</i>				
Firing threshold (mV)	V_{th}	0.8	0.8	1.35
Membrane time constant (ms)	τ_m	7.04	7.04	7.04
Membrane capacitance (pF)	C_m	3.96×10^6	3.96×10^6	3.96×10^6
Excitatory synaptic time constant (ms)	τ_{syn}^{ex}	27.0	75.0	27.0
Inhibitory synaptic time constant (ms)	τ_{syn}^{in}	50.0	50.0	50.0
Background current (pA)	I_e	4.5×10^5	9.0×10^5	0
Refractory period (ms)	t_{ref}	0	0	0
<i>Ring-attractor topology and readout</i>				
Population size	N		200	
Represented positions	N_{pos}		30	
Readout weight scale	$s_{readout}$		200	
Ridge regularization	λ		10^{-6}	
<i>Gain-modulation connectivity</i>				
Homeostatic comparator to gain	$w_{h \rightarrow g}$		1.0×10^4	
Ring to gain	$w_{r \rightarrow g}$		4.5×10^5	
Gain to ring	$w_{g \rightarrow r}$		-0.6	
Cross-inhibition	w_{cross}		-3.0×10^4	

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and introduction state the proposed ring-attractor control architecture, the SO(2) transport analysis, Fourier decoder, forward-kinematic readout, and NEST/Tiago validation. These claims match the theoretical derivations and experiments reported in the paper.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: The paper discusses target-dependent residual error, joint-dependent robot accuracy, calibration sensitivity, known-kinematics assumptions, and that the controller is not intended to replace optimized PID or model-based controllers.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: The paper provides derivations for the SO(2) reduction, gain-modulation transport equation, Fourier phase decoder, and forward-kinematic product-feature readout in the appendix. The assumptions include a stable single-bump regime, small phase shifts, approximate bump-shape preservation, and known revolute kinematics.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes].

Justification: The architecture, equations, model parameters, NEST implementation details, robot interface, and evaluation metrics are described in the main text and appendix. The Tiago experiments require a ROS/Tiago setup, but the paper reports the relevant control interface and logged-state evaluation procedure. Code will also be provided.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: An anonymized supplementary code package is provided with scripts for the NEST simulations, Fourier decoder analysis, forward-kinematic decoder analysis, and plotting from saved Tiago logs. Robot-control scripts are included for completeness, but full execution requires a compatible ROS/Tiago environment. The NEST experiments require a local NEST installation; a non-anonymous Docker image may be released after review, but is not included at submission to preserve anonymity.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: The paper specifies the NEST neuron model, ring size, represented positions, readout parameters, gain-modulation connectivity, robot interface, evaluation metrics, and appendix derivations.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: The paper reports repeated trials for the spiking readout and ring-controller experiments, including mean and standard deviation where applicable. Some robot results are reported as batch-wise errors rather than formal significance tests, which is appropriate for the small-scale hardware validation presented here.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [No]

Justification: The paper specifies the simulation platform and model parameters, but it does not currently report detailed compute resources such as CPU/GPU type, memory, runtime per experiment, or total compute. The NEST experiments are lightweight relative to large-scale training, while the robot experiments require a ROS/Tiago setup.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The work studies a neuromorphic control architecture using simulation and robot experiments, without human subjects, sensitive personal data, or high-risk deployed decision systems. The submission is anonymized for review.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [No]

Justification: The paper discusses possible positive scientific and robotic-control implications, but it does not include a dedicated discussion of negative societal impacts. The work is foundational neuromorphic control research and is not directly tied to a deployed application.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper does not release high-risk pretrained models, generative models, scraped datasets, or other assets with a high risk of misuse.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.

- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [No]

Justification: Existing tools and platforms such as NEST, ROS, and Tiago are cited or named, but the current paper does not explicitly list their licenses or terms of use. These should be documented in the supplementary code README if code is released.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: The submission includes anonymized supplementary code for reproducibility. The documentation specifies which components can be run as standalone NEST and analysis scripts, and which robot-control components require a compatible ROS/Tiago environment and are therefore provided for completeness rather than direct execution by all reviewers.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The paper does not involve crowdsourcing, user studies, or research with human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: The paper does not involve human subjects, so IRB or equivalent approval is not applicable.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A]

Justification: LLMs are not used as part of the proposed method, experiments, analysis, or scientific contribution.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.